

New Objective Functions for Social Collaborative Filtering

Joseph Noel
ANU; Canberra, AU
jinonoel@gmail.com

Peter Christen
ANU; Canberra, AU
first.last@anu.edu.au

Scott Sanner
NICTA & ANU; Canberra, AU
first.last@nicta.com.au

Lexing Xie
ANU; Canberra, AU
first.last@anu.edu.au

Khoi-Nguyen Tran
ANU; Canberra, AU
first.last@anu.edu.au

Edwin Bonilla
NICTA & ANU; Canberra, AU
first.last@nicta.com.au

ABSTRACT

This paper examines the problem of social *collaborative filtering* (CF) algorithms to recommend items of interest to users in a social network setting. Unlike standard CF algorithms using relatively simple user and item features, recommendation in social networks poses the more complex problem of learning user preferences from a rich and complex set of user profile and interaction information. Many existing *social CF* methods have extended traditional CF *matrix factorization*, but have overlooked important aspects germane to the social setting; specifically, existing matrix factorization methods (a) do not exploit user features in all aspects of learning, (b) do not permit directly modeling user-to-user information diffusion, and (c) use objectives that treat users as globally (dis)similar even though they may only be (dis)similar in specific latent areas of interest. This paper proposes a unified framework for social CF matrix factorization that addresses (a)–(c) by introducing novel objective functions for training. We demonstrate that optimizing these new objectives significantly outperforms a variety of CF and social CF baselines on live user trials in a custom-developed Facebook App involving data collected over two months from over 100 App users and their 34,000+ friends.

Categories and Subject Descriptors

H.3.3 [Information Search and Retrieval]: information filtering

General Terms

Algorithms, Experimentation

Keywords

social networks, collaborative filtering, machine learning

1. INTRODUCTION

Given the vast amount of content available on the Internet, finding information of personal interest (news, blogs,

videos, movies, books, etc.) is often like finding a needle in a haystack. Recommender systems based on *collaborative filtering* (CF) [15] aim to address this problem by leveraging the preferences of a user population under the assumption that similar users will have similar preferences.

As the web has become more social with the emergence of Facebook, Twitter, LinkedIn, and most recently Google+, this adds myriad new dimensions to the recommendation problem by making available a rich labeled graph structure of social content from which user preferences can be learned and new recommendations can be made. In this socially connected setting, no longer are web users simply described by an IP address (with perhaps associated geographical information and browsing history), but rather they are described by a rich user profile (age, gender, location, educational and work history, preferences, etc.) and a rich history of user interactions with their friends (comments/posts, clicks of the like button, tagging in photos, mutual group memberships, etc.). This rich information poses both an amazing opportunity and a daunting challenge for machine learning methods applied to social recommendation — how do we fully exploit rich social network content in recommendation algorithms?

Many existing *social CF* (SCF) approaches [10, 11, 18, 6, 12, 8] extend *matrix factorization* (MF) techniques such as [16] used in the non-social CF setting. These MF approaches have proved quite powerful and indeed, we will show empirically in Section 5 that existing social extensions of MF outperform a variety of other non-MF SCF baselines. The power of CF MF methods stems from their ability to project users and items into latent vector spaces of reduced dimensionality where each is effectively grouped by similarity; in turn, the power of many of the SCF MF extensions stems from their ability to use social network evidence to further constrain (or regularize) latent user projections.

Given the strong performance of existing MF approaches to SCF, we aim to further improve on their performance in this paper. To do this, we have first identified a number of major deficiencies of existing SCF MF objective functions:

- (a) **Feature-based user similarity learning:** Existing SCF MF objectives do not fully exploit user features when learning user similarity based on observed interactions. For example, one might enforce that two users are similar when their gender matches, but existing SCF MF objectives do not allow *learning* such a property in cases where the data supports it.
- (b) **Direct learning of user-to-user information diffusion:** Existing SCF MF objectives do not permit directly modeling user-to-user information diffusion ac-

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

WWW '12 Lyon, France

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$10.00.

cording to the social graph structure. For example, if a certain user *always* likes content liked by a friend, this cannot be directly *learned* by optimizing existing SCF MF objectives.

- (c) **Learning restricted interests:** Existing SCF MF objectives treat users as globally (dis)similar although they may only be (dis)similar in specific latent areas of interest. For example, a friend and their co-worker may both like technology-oriented news content, but have differing interests when it comes to politically-oriented news content. Existing SCF MF objective functions do not leverage *co-preference (dis)agreement* (i.e., whether two users agreed or disagreed in their rating of the *same* item) to encourage *learning* restricted latent areas where two users are (dis)similar.

This paper addresses (a)–(c) by proposing novel objective functions in a unified latent factorization framework for SCF. We present results of our proposed recommendation algorithms in online human trials of a custom-developed Facebook App for link recommendation that involves data collected over three months from over 100 active App users and their 34,000+ friends. These results show that the extensions proposed to resolve (a)–(c) outperform a wide range of previously existing SCF algorithms.

To this end, the remaining sections of the paper are organized as follows:

- **Section 2 – Definitions and Background:** We define mathematical notation used throughout the paper along with a discussion of related work and all baseline comparison methods we use in this paper.
- **Section 3 – New Objective Functions for SCF:** We propose novel extensions to address (a)–(c) within a unified mathematical framework. Gradients for optimizing these functions are provided in Appendix A.
- **Section 4 – Evaluation Framework:** We discuss the details of our Facebook App for link recommendation as well as our evaluation methodology for online (live user trial) experimentation.
- **Section 5 – Empirical Results:** In our first online trial, we identify an MF-based SCF method as the best-performing baseline algorithm. This directly motivated our decision to extend MF-based methods in Section 3 and hence we compare these new objectives to the best-performing SCF MF baseline from the first trial. After showing where each extension improves on the baseline, we propose a number of additional experiments to better understand trends and patterns in our social recommendation setting such as the use of explicit like/dislike feedback vs. link click information, the impact of textual link context, and the relationship between like popularity and user preference.
- **Section 6 – Conclusions:** We summarize our conclusions from the empirical evaluations and outline future research directions stemming from the current work.

All combined, this paper represents a critical step forward in powerful SCF recommendation algorithms based on latent factorization methods and their ability to fully exploit the breadth of information available on social networks.

2. DEFINITIONS AND BACKGROUND

Collaborative filtering (CF) [15] is the task of predicting whether, or how much, a user will like (or dislike) an item by leveraging knowledge of that user’s preferences *as well as those of other users*. While collaborative filtering need not take advantage of user or item features (if available), a separate approach of *content-based filtering* (CBF) [7] makes individual recommendations by generalizing from the item features of those items the user has explicitly liked or disliked. What distinguishes CBF from CF is that CBF requires item features to generalize whereas CF requires multiple users to generalize; however, CBF and CF are not mutually exclusive and recommendation systems often combine the two approaches [1]. When a CF method makes use of item and user features as well as multiple users, we refer to it as CF although in some sense it may be viewed as a combined CF and CBF approach.

We define *social CF* (SCF) simply as the task of CF augmented with additional social network information such as the following:

- Expressive personal profile content: gender, age, places lived, schools attended; favorite books, movies, quotes; online photo albums (and associated comment text).
- Explicit friendship or trust relationships.
- Content that users have personally posted (often text, images, and links).
- Content of public (and if available, private) interactions between users (often text, images and links).
- Evidence of external interactions between users such as being jointly tagged in photos or videos.
- Expressed preferences (likes/dislikes of posts and links).
- Group memberships (often for hobbies, activities, social or political discussion).

We note that CF is possible in a social setting without taking advantage of the above social information, hence we include CF baselines in our later experiments on SCF.

2.1 Notation

We attempt to present all algorithms for CF and SCF using the following mathematical notation:

- N users. For methods that can exploit user features, we define an I -element user feature vector $\mathbf{x} \in \mathbb{R}^I$ (alternately if a second user is needed, $\mathbf{z} \in \mathbb{R}^I$). For methods that do not use user feature vectors, we simply assume \mathbf{x} is an index $\mathbf{x} \in \{1 \dots N\}$ and that $I = N$.
- M items. For methods that can exploit item features, we define a J -element feature vector $\mathbf{y} \in \mathbb{R}^J$. The feature vectors for users and items can consist of any real-valued features as well as $\{0, 1\}$ features like user and item IDs. For methods that do not use item feature vectors, we simply assume \mathbf{y} is an index $\mathbf{y} \in \{1 \dots M\}$ and that $J = M$.
- A (non-exhaustive) data set D of single user *preferences* of the form $D = \{(\mathbf{x}, \mathbf{y}) \rightarrow R_{\mathbf{x}, \mathbf{y}}\}$ where the binary *response* $R_{\mathbf{x}, \mathbf{y}} \in \{0 \text{ (dislike)}, 1 \text{ (like)}\}$.

- A (non-exhaustive) data set C of *co-preferences* (cases where *both* users \mathbf{x} and \mathbf{z} expressed a preference for \mathbf{y} – not necessarily in agreement) derived from D of the form $C = \{(\mathbf{x}, \mathbf{z}, \mathbf{y}) \rightarrow P_{\mathbf{x}, \mathbf{z}, \mathbf{y}}\}$ where co-preference class $P_{\mathbf{x}, \mathbf{z}, \mathbf{y}} \in \{-1$ (disagree), 1 (agree) $\}$. Intuitively, if *both* user \mathbf{x} and \mathbf{z} liked or disliked item \mathbf{y} then we say they *agree*, otherwise if one liked the item and the other disliked it, we say they *disagree*.
- A similarity rating $S_{\mathbf{x}, \mathbf{z}}$ between any users \mathbf{x} and \mathbf{z} . This is used to summarize all social interaction between user \mathbf{x} and user \mathbf{z} in the term $S_{\mathbf{x}, \mathbf{z}} \in \mathbb{R}$. A definition of $S_{\mathbf{x}, \mathbf{z}} \in \mathbb{R}$ that has been useful is the following average-normalized measure of user interactions:

$$Int_{\mathbf{x}, \mathbf{z}} = \frac{\# \text{ interactions between } \mathbf{x} \text{ and } \mathbf{z}}{\frac{1}{N^2} \sum_{\mathbf{x}', \mathbf{z}'} \# \text{ interactions between } \mathbf{x}' \text{ and } \mathbf{z}'}$$

$$S_{\mathbf{x}, \mathbf{z}} = \ln(Int_{\mathbf{x}, \mathbf{z}}) \quad (1)$$

How “# interactions between \mathbf{x} and \mathbf{z} ” is explicitly defined is specific to a social network setting and hence we defer details of the particular method user for evaluations in this paper to Section 4.2.4.

We also define $S_{\mathbf{x}, \mathbf{z}}^+$, a *non-negative* variant of $S_{\mathbf{x}, \mathbf{z}}$:

$$S_{\mathbf{x}, \mathbf{z}}^+ = \ln(1 + Int_{\mathbf{x}, \mathbf{z}}) \quad (2)$$

- A set $friends_{\mathbf{x}}$ such that $\mathbf{z} \in friends_{\mathbf{x}}$ iff \mathbf{z} is officially denoted as a *friend* of \mathbf{x} on the social network.

Having now defined notation, we proceed to survey a number of CBF, CF, and SCF algorithms including all of those compared to or extended in this paper.

2.2 Content-based Filtering (CBF)

Since our objective in this work is to classify whether a user likes an item or not (i.e., a binary objective), we focus on binary classification-based CBF approaches in this paper. While a variety of classifiers may work well, we choose the *support vector machine* (SVM) [5] since it is well-known for its state-of-the-art classification performance.

For the experiments in this paper, we use a *linear* SVM (implemented in the *LibSVM* [4] toolkit) with feature vector $\mathbf{f} \in \mathbb{R}^F$ derived from $(\mathbf{x}, \mathbf{y}) \in D$, denoted as $\mathbf{f}_{\mathbf{x}, \mathbf{y}}$. A linear SVM learns a weight vector $\mathbf{w} \in \mathbb{R}^F$ such that $\mathbf{w}^T \mathbf{f}_{\mathbf{x}, \mathbf{y}} > 0$ indicates a like (1) classification of $\mathbf{f}_{\mathbf{x}, \mathbf{y}}$ and $\mathbf{w}^T \mathbf{f}_{\mathbf{x}, \mathbf{y}} \leq 0$ indicates a dislike (0) classification.

A detailed list of features used in the SVM for the Facebook link recommendation task evaluated in this paper are defined in Section 4.2 — these include *user features* such as age and gender (binary) and *item features* such as popularity (number of times the item was shared). Going one step beyond standard CBF, our SVM features also include *joint* user and item features from the social network, in particular binary *information diffusion* [3] features for *each* friend $\mathbf{z} \in friends_{\mathbf{x}}$ indicating if \mathbf{z} liked (or disliked) \mathbf{y} . Crucially we note that our SVM implementation of CBF using social network features actually represents a *social CBF* extension since it can learn when a friend \mathbf{z} ’s preference for items are predictive of user \mathbf{x} ’s preferences.

2.3 Collaborative Filtering (CF)

2.3.1 *k*-Nearest Neighbor

One of the most common forms of CF is the nearest neighbor approach [2]. There are two main variants of nearest neighbors for CF, *user-based* and *item-based* — both methods generally assume that no user or item features are provided, so here \mathbf{x} and \mathbf{y} are simply respective user and item indices. When the number of users is far fewer than the number of items, it has been found that the user-based approach usually provides better predictions as well as being more efficient in computations [2]; this holds for the evaluation in this paper, so we focus on the user-based approach.

Given a user \mathbf{x} and an item \mathbf{y} , let $N(\mathbf{x} : \mathbf{y})$ be the set of *user* nearest neighbors of \mathbf{x} that have also given a rating for \mathbf{y} and let $S_{\mathbf{x}, \mathbf{z}}$ be the cosine similarity (i.e., normalized dot product) between two vectors of ratings for users \mathbf{x} and \mathbf{z} (when both have rated the same item). Following [2], the predicted rating $\hat{R}_{\mathbf{x}, \mathbf{y}} \in [0, 1]$ that the user \mathbf{x} gives item \mathbf{y} can then be calculated as

$$\hat{R}_{\mathbf{x}, \mathbf{y}} = \frac{\sum_{\mathbf{z} \in N(\mathbf{x} : \mathbf{y})} S_{\mathbf{x}, \mathbf{z}} R_{\mathbf{z}, \mathbf{y}}}{\sum_{\mathbf{z} \in N(\mathbf{x} : \mathbf{y})} S_{\mathbf{x}, \mathbf{z}}} \quad (3)$$

2.3.2 Matrix Factorization (MF) Models

Another common approach to CF attempts to factorize an (incomplete) matrix R of dimension $I \times J$ containing observed ratings $R_{\mathbf{x}, \mathbf{y}}$ (note that \mathbf{x} and \mathbf{y} are assumed to row and column indices) into a product $R \approx U^T V$ of real-valued rank- K matrices U and V :

$$U = \begin{bmatrix} U_{1,1} & \cdots & U_{1,I} \\ \vdots & U_{k,i} & \vdots \\ U_{K,1} & \cdots & U_{K,I} \end{bmatrix} \quad V = \begin{bmatrix} V_{1,1} & \cdots & V_{1,J} \\ \vdots & V_{k,j} & \vdots \\ V_{K,1} & \cdots & V_{K,J} \end{bmatrix}$$

In this setting, we *do not* leverage user and item features; hence, the \mathbf{x} and \mathbf{y} indices simply pick out the respective rows and columns of U and V such that $U_{\mathbf{x}}^T V_{\mathbf{y}}$ acts as a measure of affinity between user \mathbf{x} and item \mathbf{y} in respective K -dimensional latent spaces $U_{\mathbf{x}}$ and $V_{\mathbf{y}}$.

However, there remains the question of how we can learn U and V given that R is incomplete (i.e., it contains missing entries since D is generally non-exhaustive). The answer is simple: we need only define a reconstruction error objective we wish to minimize as a function of U and V and then use gradient descent to optimize it; formally then, we can optimize the following MF objective [16]:

$$\sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} (R_{\mathbf{x}, \mathbf{y}} - U_{\mathbf{x}}^T V_{\mathbf{y}})^2 \quad (4)$$

While this objective is technically bilinear, we can easily apply an *alternating gradient descent* approach to approximately optimize it and determine good projections U and V that (locally) minimize the reconstruction error of the observed responses $R_{\mathbf{x}, \mathbf{y}}$ [16].

2.3.3 Social Collaborative Filtering (SCF)

In this work on *social CF* (SCF), we opt to focus on extending MF-based SCF objectives since gradient descent MF optimization approaches allow us a flexible language to design and optimize objective functions that take into account a vast array of social network information.

To date, there are essentially two general classes of MF methods applied to SCF of which the authors are aware. The first class of social MF methods can be termed as *social*

regularization approaches in that they somehow constrain the latent projection of users according to social network information. There are two closely related social regularization methods that directly constrain $U_{\mathbf{x}}$ and $U_{\mathbf{z}}$ for user \mathbf{x} and \mathbf{z} based on evidence $S_{\mathbf{x},\mathbf{z}}$ of interaction between \mathbf{x} and \mathbf{z} . The first class of methods are simply termed *social regularization* [18, 6] where $\langle \cdot, \cdot \rangle$ denotes an inner product

$$\sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}_{\mathbf{x}}} \frac{1}{2} (S_{\mathbf{x},\mathbf{z}} - \langle U_{\mathbf{x}}, U_{\mathbf{z}} \rangle)^2 \quad (5)$$

while the second class of methods are termed *social spectral regularization* [12, 8]

$$\sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}_{\mathbf{x}}} \frac{1}{2} S_{\mathbf{x},\mathbf{z}}^+ \|U_{\mathbf{x}} - U_{\mathbf{z}}\|_2^2 \quad (6)$$

We refer to the latter as *spectral regularization* methods since they are identical to the objectives used in spectral clustering [13]. The idea behind both variants of social regularization should be apparent: the larger $S_{\mathbf{x},\mathbf{z}}$ or $S_{\mathbf{x},\mathbf{z}}^+$, the more $U_{\mathbf{x}}$ and $U_{\mathbf{z}}$ need to be similar (according to slightly different metrics) in order to minimize the given objective.

The *SoRec* system [11] proposes a slight twist on social spectral regularization in that it learns a third $N \times N$ (n.b., $I = N$) *interactions matrix* Z , and uses $U_{\mathbf{z}}^T Z_{\mathbf{z}}$ to predict user-user interaction preferences in the same way that standard CF uses V in $U_{\mathbf{x}}^T V_{\mathbf{y}}$ to predict user-item ratings. *SoRec* also uses a sigmoidal transform $\sigma(o) = \frac{1}{1+e^{-o}}$ since $\bar{S}_{\mathbf{x},\mathbf{z}}$ is $S_{\mathbf{x},\mathbf{z}}$ restricted to the range $[0, 1]$ (e.g., $\bar{S}_{\mathbf{x},\mathbf{z}} = \sigma(S_{\mathbf{x},\mathbf{z}})$):

$$\sum_{\mathbf{z}} \sum_{\mathbf{z} \in \text{friends}_{\mathbf{x}}} \frac{1}{2} (\bar{S}_{\mathbf{x},\mathbf{z}} - \sigma(\langle U_{\mathbf{x}}, Z_{\mathbf{z}} \rangle))^2 \quad (7)$$

The second class of SCF MF approaches represented by the single exemplar of the *Social Trust Ensemble* (STE) [10] can be termed as a *weighted friend average* approach since this approach simply composes a prediction for item \mathbf{y} from an α -weighted average ($\alpha \in [0, 1]$) of a user \mathbf{x} 's predictions *as well as* their friends (\mathbf{z}) predictions (as evidenced by the additional $\sum_{\mathbf{z}}$ in the objective below):

$$\sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} (R_{\mathbf{x},\mathbf{y}} - \sigma(\alpha U_{\mathbf{x}}^T V_{\mathbf{y}} + (1 - \alpha) \sum_{\mathbf{z} \in \text{friends}_{\mathbf{x}}} U_{\mathbf{x}}^T V_{\mathbf{z}}))^2 \quad (8)$$

As for the MF CF methods, all MF SCF methods can be optimized by alternating gradient descent on the respective matrix parameterizations; we refer the interested reader to each paper for further details.

3. NEW OBJECTIVE FUNCTIONS FOR SCF

Having surveyed previous CF work, especially MF-based CF and SCF methods, we now present the major conceptual contributions of the paper. However, before we get into details, we introduce a unified matrix factorization framework for optimizing all MF objectives evaluated in this paper, including newly proposed objectives to address observed deficiencies of SCF MF methods as outlined in Section 1.

3.1 A Composable Objective Framework

We take a composable approach to MF-based (S)CF, where an optimization objective *Obj* is composed of sums of one

or more objective components:

$$\text{Obj} = \sum_i \lambda_i \text{Obj}_i \quad (9)$$

Because each objective may be weighted differently, a weighting term $\lambda_i \in \mathbb{R}$ is included for each component. In the current work, we manually tuned each λ_i , except for the last i in \sum_i , which can be set as $\lambda_i = 1$ without loss of generality.

Most target predictions in this paper are binary ($\{0, 1\}$), therefore a sigmoidal transform $\sigma(o) = \frac{1}{1+e^{-o}}$ of a prediction $o \in \mathbb{R}$ may be used to squash it to the range $[0, 1]$. Where the σ transform may be optionally included, this is written as $[\sigma]$. While σ transforms are generally advocated for real-valued regressor outputs when used in a classification setting, we note that our experiments showed little variation in results whether including or omitting it, although including it tended to slow the convergence of gradient optimization. Nonetheless, where appropriate, we include the possibility of a σ transform since it may prove useful in other settings.

3.2 Existing Objective Functions

For completeness, we first cover a number of known objective components that are used in the objectives evaluated and extended in this paper. A discussion of gradient optimization along with all necessary derivatives for these objectives is provided in Appendix A.

3.2.1 Matchbox-style Matrix Factorization (Obj_{pmcf})

In Section 2.3.2, we discussed an MF objective (4) that *did not* make use of user or item features. However, if we *do* have user feature vector \mathbf{x} and item feature vector \mathbf{y} , we can respectively represent the latent projections of user and item as $(U_{\mathbf{x}})_{1 \dots K}$ and $(V_{\mathbf{y}})_{1 \dots K}$ and hence use $\langle U_{\mathbf{x}}, V_{\mathbf{y}} \rangle = \mathbf{x}^T U^T V \mathbf{y}$ as a measure of affinity between user \mathbf{x} and item \mathbf{y} . Substituting the feature-based $\mathbf{x}^T U^T V \mathbf{y}$ for the featureless $U_{\mathbf{x}}^T V_{\mathbf{y}}$ in (4), we arrive at the form of the basic CF objective function used in *Matchbox* [17] — although *Matchbox* used Bayesian optimization methods, we can easily express its objective in the following log likelihood form:

$$\text{Obj}_{pmcf} = \sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} (R_{\mathbf{x},\mathbf{y}} - [\sigma] \mathbf{x}^T U^T V \mathbf{y})^2 \quad (10)$$

3.2.2 Regularization of U , V , & \mathbf{w} (Obj_{ru} , Obj_{rv} , $\text{Obj}_{\mathbf{w}}$)

To help in generalization, it is important to regularize any free matrix parameters U and V (e.g., from Section 3.2.1) or vector parameters \mathbf{w} (e.g., from Section 2.2) to prevent overfitting in the presence of sparse data. This can be done with a simple L_2 regularizer that models a spherical Gaussian prior of 0 on the parameters. This regularization component can be specified for U , V , and \mathbf{w} simply as follows:

$$\begin{aligned} \text{Obj}_{ru} &= \frac{1}{2} \|U\|_{\text{Fro}}^2 = \frac{1}{2} \text{tr}(U^T U) & \text{Obj}_{rv} &= \frac{1}{2} \text{tr}(V^T V) \\ \text{Obj}_{\mathbf{w}} &= \frac{1}{2} \|\mathbf{w}\|_2^2 = \frac{1}{2} \mathbf{w}^T \mathbf{w} \end{aligned} \quad (11)$$

3.3 New Objective Functions

Now we return to our observations concerning the deficiencies of existing SCF MF methods as outlined in Section 1 and propose new objective functions to address these issues. Gradient optimization for these new objectives and all necessary derivatives are covered in Appendix A.

3.3.1 Feature Social Regularization (Obj_{rs} & Obj_{rss})

Our previous discussion of SCF methods in Section 2.3.3 covered three different methods for *social regularization* — that is, constraining users to be similar based on evidence from the social network. However, none of these previous three SCF social regularization methods exploited user features in the *learning* process; more precisely $U_{\mathbf{x}}$ and $U_{\mathbf{z}}$ were regularized, but not the feature-based latent spaces $U\mathbf{x}$ and $U\mathbf{z}$. Hence if a gender difference in \mathbf{x} and \mathbf{z} was the crucial factor to differentiating the latent spaces of each user, it *could* be learned if we had a way of socially regularizing $U\mathbf{x}$ and $U\mathbf{z}$ directly. To address this, we provide two variants of *feature-based social regularization*.

The first new objective is an extension of simple *social regularization* [18, 6] by incorporating user features into that objective:

$$\begin{aligned} Obj_{rs} &= \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \frac{1}{2} (S_{\mathbf{x},\mathbf{z}} - \langle U\mathbf{x}, U\mathbf{z} \rangle)^2 \\ &= \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \frac{1}{2} (S_{\mathbf{x},\mathbf{z}} - \mathbf{x}^T U^T U \mathbf{z})^2 \end{aligned} \quad (12)$$

Alternately, we could extend *social spectral regularization* [12, 8] by incorporating user features into the objective:

$$\begin{aligned} Obj_{rss} &= \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \frac{1}{2} S_{\mathbf{x},\mathbf{z}}^+ \|U\mathbf{x} - U\mathbf{z}\|_2^2 \\ &= \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \frac{1}{2} S_{\mathbf{x},\mathbf{z}}^+ (\mathbf{x} - \mathbf{z})^T U^T U (\mathbf{x} - \mathbf{z}) \end{aligned} \quad (13)$$

While we note these extensions are straightforward, they have the simple property that they allow the system to learn the latent user projection matrix U as a function of user features in order to minimize the social regularization penalty. Compared to non-socially regularized CF objectives, we will demonstrate that both of these approaches are quite powerful in Section 5.

3.3.2 Hybrid Information Diffusion + SCF (Obj_{phy})

One major weakness of MF methods is that they cannot model direct joint features over user and items — they must model user and item features independently in order to compute the independent latent projections $U\mathbf{x}$ and $U\mathbf{z}$. Unfortunately, this prevents standard MF objectives from modeling direct user-to-user information diffusion [3] — the unidirectional flow of links from one user to another. This is problematic because if user \mathbf{x} *always* likes what \mathbf{z} has posted or liked, then we would like to shortcut the latent representation and simply learn to recommend user \mathbf{z} ’s likes and posts to user \mathbf{x} .

We fix this deficit of MF by introducing another objective component in addition to the standard MF objective, which serves as a simple linear regressor for such information diffusion observations. The resulting hybrid objective component then becomes a combination of latent MF and linear regression objectives.

For the linear regressor $\mathbf{w}^T \mathbf{f}_{\mathbf{x},\mathbf{y}}$, we make use of the *same* weight vector \mathbf{w} and feature vector $\mathbf{f}_{\mathbf{x},\mathbf{y}}$ mentioned in Section 2.2; $\mathbf{f}_{\mathbf{x},\mathbf{y}}$ is fully defined for our empirical evaluation in Section 4.2. As previously noted, $\mathbf{f}_{\mathbf{x},\mathbf{y}}$ includes *joint* user and item features from the social network, in particular binary *information diffusion* [3] features for *each* friend $\mathbf{z} \in$

$\text{friends}_{\mathbf{x}}$ indicating if \mathbf{z} liked (or disliked) \mathbf{y} . As a consequence, learning \mathbf{w} allows the linear regressor to predict in a personalized way for any user \mathbf{x} whether they are likely to follow their friend \mathbf{z} ’s preference for \mathbf{y} .

Formally, to define our hybrid information diffusion + SCF objective, we combine the output of the linear regression prediction with the Matchbox prediction, to get a hybrid objective component. The full objective function for this hybrid model is then

$$Obj_{phy} = \sum_{(\mathbf{x},\mathbf{y}) \in D} \frac{1}{2} (R_{\mathbf{x},\mathbf{y}} - [\sigma] \mathbf{w}^T \mathbf{f}_{\mathbf{x},\mathbf{y}} - [\sigma] \mathbf{x}^T U^T V \mathbf{y})^2 \quad (14)$$

While again we note that this simple hybrid MF + linear regression objective is straightforward, the ability to use joint user and item features to model information diffusion between users turns out to be extremely powerful in our later experiments.

3.3.3 Co-preference Regularization (Obj_{rsc})

A crucial aspect missing from other SCF methods is that while two users may not be globally similar or opposite in their preferences, there may be sub-areas of their interests which can be correlated to each other. For example, two friends may have similar interests concerning technology news, but different interests concerning political news. *Co-preference regularization* aims to learn such selective co-preferences. The motivation is to constrain users \mathbf{x} and \mathbf{z} who have similar or opposing preferences to be similar or opposite in the same latent space relevant to item \mathbf{y} .

We use $\langle \cdot, \cdot \rangle_{\bullet}$ to denote a re-weighted inner product. The purpose of this inner product is to tailor the latent space similarities or dissimilarities between users to specific sets of items. This addresses the issue detailed in the previous paragraph by allowing users \mathbf{x} and \mathbf{z} to be similar or opposite in the same latent space relevant only to item \mathbf{y} .

The objective component for *Co-preference regularization* along with its expanded form is

$$\begin{aligned} Obj_{rsc} &= \sum_{(\mathbf{x},\mathbf{z},\mathbf{y}) \in C} \frac{1}{2} (P_{\mathbf{x},\mathbf{z},\mathbf{y}} - \langle U\mathbf{x}, U\mathbf{z} \rangle_{V\mathbf{y}})^2 \\ &= \sum_{(\mathbf{x},\mathbf{z},\mathbf{y}) \in C} \frac{1}{2} (P_{\mathbf{x},\mathbf{z},\mathbf{y}} - \mathbf{x}^T U^T \text{diag}(V\mathbf{y}) U \mathbf{z})^2 \end{aligned} \quad (15)$$

We might also define a *social co-preference spectral regularization* approach, but our experiments so far have not indicated this works as well as the above objective.

In contrast to social regularization defined previously, co-preference regularization does not require knowledge of friendships or user interactions to determine co-preferences and hence can find correlations between users who are not friends — this important observation will indeed surface in our forthcoming experimental evaluation.

4. EVALUATION FRAMEWORK

In this section we discuss our Facebook Link Recommendation (LinkR) application, definitions for notation in the context of this application, and our evaluation methodology.

4.1 Link Recommendation App on Facebook

To evaluate existing and newly proposed (S)CF methods discussed in this paper, we created a Facebook application



Figure 1: The Facebook LinkR App showing two link recommendations to a user. One link recommendation is from a friend and includes the friend’s commentary on the link as well as the link description. The other link recommendation is from a non-friend and hence only shows the link description. Users have the option of liking or disliking each recommendation as well as providing explicit feedback.

(i.e., a Facebook “App”) that recommends links to users everyday, where the users may give their feedback on the links indicating whether they *liked* it or *disliked* it. Figure 1 shows our Facebook LinkR App as it appears to users.

The functionalities of the LinkR application on a daily basis are as follows:

1. Collect data that have been shared by users and their friends on Facebook.
2. Initiate retraining of all active (S)CF link recommendation algorithms on the latest collected data.
3. Following retraining, recommend three links to the users according to their assigned recommendation algorithm.
4. Collect feedback from the users on whether they liked or disliked the recommendations as well as any additional commentary the user wishes to provide.

Details of the (S)CF link recommendation algorithms and user assignments will be discussed shortly, but first we cover the specific data collected by the LinkR App and made available for use by the recommendation algorithms.

4.2 Facebook Data Collected

At its peak membership, 108 users had elected to install the Facebook App developed for this project. From this

user base, we were able to gather data on 34,860 users and 437,023 links in total by the end of the evaluation period.

4.2.1 User Data

Data that are collected and used to define the user feature vector \mathbf{x} introduced in Section 2.1 for the LinkR Facebook App are defined as follows:

- $\forall id, \mathbf{x}_{id} = id \in \{0, 1\}$: every unique Facebook ID (user) recorded in the App was assigned its own binary indicator in \mathbf{x} ; all such indicators are enforced to be mutually exclusive.
- $gender \in \{0 \text{ female}, 1 \text{ male}\}$.
- $age \in \mathbb{N}$.

We note that the indicator of friendships for \mathbf{x} is stored in the $friends_{\mathbf{x}}$ set defined in Section 2.1 and used in various previous objective definitions, but not explicitly stored in \mathbf{x} .

4.2.2 Link Data

Data that are collected and used to define the item feature vector \mathbf{y} introduced in Section 2.1 for the LinkR Facebook App are defined as follows:

- $\forall id, \mathbf{y}_{poster} = id \in \{0, 1\}$: binary indicator feature for the id of the user who posted the link; all such binary indicator features are mutually exclusive.
- $\forall id, \mathbf{y}_{wall} = id \in \{0, 1\}$: binary indicator feature for the id of the user on whose wall the link was posted; all such binary indicator features are mutually exclusive.
- Count of total link “likes” on Facebook.
- Count of total link shares on Facebook.
- Count of total link comments posted on Facebook.

4.2.3 Joint User and Link Data

The feature vector $\mathbf{f}_{\mathbf{x}, \mathbf{y}}$ used in Sections 2.2 and 3.3.2 for the LinkR Facebook App is defined as the *concatenation* of \mathbf{x} , \mathbf{y} (above) and the following additional social network information diffusion features:

- $\forall \mathbf{z} \in friends_{\mathbf{x}}, \mathbf{z} \text{ liked } \mathbf{x} \in \{0, 1\}$: for every friend \mathbf{z} of user \mathbf{x} , we have a binary information diffusion feature indicating whether user \mathbf{z} liked item \mathbf{y} (recall that $\mathbf{f}_{\mathbf{x}, \mathbf{y}}$ is built w.r.t. a specific user \mathbf{x} and item \mathbf{y}).

4.2.4 Interaction Data

The Facebook social network interactions between users \mathbf{x} and \mathbf{z} that we count (equally weighted) to define $\# \text{ interactions between } \mathbf{x} \text{ and } \mathbf{z}$ in Section 2.1 are defined as follows:

1. Being friends.
2. Posting, liking, or commenting on an item (link, photo, video, photo, or message) on a user’s wall.
3. Being tagged together in the same photo or video.
4. Tagging themselves as attending the same school or class, playing sports together, working together for the same company or on the same project.

4.3 Live Online Recommendation Trials

For the recommendations made to the LinkR application users, we select only links posted in the most recent two weeks that the user has not already liked (or disliked). We use only the links from the last two weeks since an informal user study has indicated a preference for recent links. Indeed, older links have a greater chance of being outdated and are also likely to represent broken links that are no longer accessible. We have chosen to recommend three links per day with the aim of avoiding position bias and information overload that may occur with longer recommendation lists.

For the live trials, Facebook users who installed the LinkR application were *randomly assigned one of four algorithms in each of two trials*. Specific algorithms trialed will be discussed in Section 5. LinkR users were not informed which algorithm was assigned to them. As demonstrated in Figure 1, we distinguish our recommended links into two major classes: (1) links that were posted by the LinkR user’s friends, and (2) links that were posted by users other than the LinkR user’s friends. LinkR users were encouraged to rate each link as like or dislike. In turn these ratings became part of the training data for the recommendation algorithms, and thus were used to improve the performance of the algorithms over time. LinkR users were allowed to provide feedback comments on specific links if they wished; based on repeated user commentary from the first trial period, we filtered out non-English links and links lacking any textual descriptions from the recommendations made on the second trial in order to prevent user annoyance.

5. EMPIRICAL RESULTS

In this section we discuss the algorithms evaluated in our two online trials¹ and additional analysis regarding trends and patterns in our social recommendation setting.

5.1 First Trial

In our first trial, our objective was to evaluate four CF and SCF approaches to establish which was the most promising direction for SCF extension:

1. ***k*-Nearest Neighbor (KNN)**: See Section 2.3.1.
2. **Support Vector Machines (SVM)**: See Section 2.2.
3. **Matchbox (Mbox)**: Gradient optimization of the L_2 regularized Matchbox objective:

$$Obj_{pmcf} + \lambda Obj_{ru} + \lambda Obj_{rv}$$

4. **Social Matchbox (Soc. Mbox)**: Gradient optimization of the L_2 and *feature socially regularized* Matchbox objective:

$$Obj_{pmcf} + \lambda_{rs} Obj_{rs} + \lambda Obj_{ru} + \lambda Obj_{rv}$$

Mbox and Soc. Mbox are optimized via gradient descent as outlined in Appendix ???. KNN and MBox may be viewed as pure CF methods. As discussed previously, both SVM and

¹All code used in these experiments is available at <http://code.google.com/p/social-recommendation/>. The conditions of our ethics approval #2011/142 from the Australian National University for conducting human trials on Facebook require our privacy policy (<http://dmm.anu.edu.au/linkr/website/pp.php>) to prohibit public sharing of data collected during these experiments.

Algorithm	Users
Social Matchbox	26
Matchbox	26
SVM	28
Nearest Neighbor	28

Table 1: Number of Users Assigned per Algorithm.

Soc. Mbox may be viewed as SCF methods. Social Matchbox uses the Social Regularization method to incorporate the social information of the FB data. SVM incorporates social information in the $\mathbf{f}_{x,y}$ features that it uses. Matchbox and Nearest Neighbors do not make use of any social information.

The first live user trial was run from August 25 to October 13. The algorithms were randomly distributed among the 106 users who installed the LinkR application. The distribution of the algorithms to the users are show in Table 1

Each user was recommended three links everyday and they were able to rate the links on whether they ‘Liked’ or ‘Disliked’ it. Results shown in Figure 2 are the percentage of Like ratings and the percentage of Dislike ratings per algorithm stacked on top of each other with the Like ratings on top.

5.1.1 Online Results

As shown in Figure 2, Social Matchbox was the best performing algorithm in the first trial and in fact was the only algorithm to get receive more like ratings than dislike ratings. This would suggest that using social information does indeed provide useful information that resulted in better link recommendations from LinkR.

We also look at the algorithms with the results split between friend links and non-friend links recommendations. Again, the results shown in Figure ?? are the percentage of Like ratings and the percentage of Dislike ratings per algorithm stacked on top of each other with the Like ratings on top. As shown in Figure ??, all four algorithms experienced a significant performance drop in the ratio of Likes to Dislikes when it came to recommending non-friend links. This suggests that aside from Liking or Disliking a link solely from the quality of the link being recommended, users are also more likely to Like a link simply because a friend had posted it and more likely to Dislike it because it was posted by a stranger.

5.1.2 Summary

At the end of the first trial, we have observed the following:

- Social Matchbox was the best performing algorithm in the live user trial in terms of percentage of likes.
- Social Matchbox received the highest user evaluation scores in the user survey at the end of the user trial.
- Of the various combinations of training and testing data in the offline passive experiment, we found that training on the UNION subset and testing on the APP-USER-ACTIVE-ALL subset best correlated with the results of the live user trial and the user survey. Training on the UNION dataset had advantages compared to training on the other data subsets, namely that it




Figure 2: Results of the online live trials. The percentage of Liked ratings are stacked on top of the percentage of Disliked ratings per algorithm. Social Matchbox was found to be the best performing of the four algorithms evaluated in the first trial. * Results of the online live trials, split between friends and non-friends. The percentage of Liked ratings are stacked on top of the percentage of Disliked ratings per algorithm. There is a significant drop in performance between recommending friend links and recommending non-friend links.**

had the large amount of information of the PASSIVE data and the explicit dislikes information of the ACTIVE data.

5.2 Second Trial

For the second online trial, we chose four algorithms again to randomly split between the LinkR application users. Social Matchbox was included again as a baseline since it was the best performing algorithm in the first trial. The distribution count of the algorithms to the users is shown in Table 2

The four SCF algorithms are:

- **Social Matchbox (Soc. Mbox)** : Matchbox MF + Social Regularization + L2 U Regularization + L2 V Regularization
- **Spectral Matchbox (Spec. Mbox)**: Matchbox MF + Social Spectral Regularization + L2 U Regularization + L2 V Regularization
- **Social Hybrid (Soc. Hybrid)**: Hybrid + Social Regularization + L2 U Regularization + L2 V Regularization + L2 w Regularization
- **Spectral Co-preference (Spec. CP)**: Matchbox MF + Social Co-preference Spectral Regularization + L2 U Regularization + L2 V Regularization

5.2.1 Online Results

The online experiments were switched to the new algorithms on October 13, 2011. For the online results reported here, since the second live trial is still currently ongoing, we took a snapshot of the data as it was on October 22, 2011. The algorithms were randomly distributed among the 103 users who still had the LinkR application installed. The distribution of the algorithms to the users are shown in Table 2

Results shown in Figure 3 are the percentage of Like ratings and the percentage of Dislike ratings per algorithm stacked on top of each other with the Like ratings on top.

First thing we note in Figure 3 is the decrease in performance for Social Matchbox, and in fact for all SCF algorithms in general. Except for Spectral Matchbox, they all received more Dislike ratings than Like ratings. What

Algorithm	Users
Social Matchbox	26
Spectral Matchbox	25
Spectral Co-preference	27
Social Hybrid	25

Table 2: Number of Users Assigned per Algorithm.

we noticed is that of the recommendations being made in the week that we switched over to the new algorithms, the majority of the links were about Steve Jobs, who had died the week previously. We believe that the redundancy and lack of variety of the links being recommended caused an increase in the Dislike ratings being given by users on the recommended links. Taking out the skewed results that follows an unusual event such as this, the relative algorithm performance was better.

We again split the results again between friend link recommendations and non-friend link recommendations, with the results shown in Figure 3 being the percentage of Like ratings and the percentage of Dislike ratings per algorithm stacked on top of each other with the Like ratings on top. As shown in Figure 3, all four algorithms experienced significant performance drop in the number of likes when it came to recommending non-friend links. This reflects the results of the first trial.

We note the following observations from the results shown in Figures 3:

- Compared to the other algorithms, Spectral Matchbox achieved the best ratio of likes to dislikes as seen, as seen in Figure 3. Combined with the results for Spectral Co-preference, Spectral social regularization in general appears to be a better way to socially regularize compared to social regularization. This comparison holds even when the results are split between friend links recommendations and non-friend links recommendations, as seen in Figure 3.
- When looking at just the friend link recommendations in Figure 3, Social Hybrid was the best performing algorithm. This result comes from the user-user information diffusion among its friends that Social Hybrid learns, which could not be learned by the other SCF




Figure 3: Results of online live trials. The percentage of Liked ratings are stacked on top of the percentage of Disliked ratings per algorithm. Spectral Matchbox achieved the highest ratio of likes to dislikes among the four algorithms. Spectral social regularization in general appears to be a better way to socially regularize compared to social regularization. *** Results of the online live trials, split between friends and non-friends. As in the first trial, there is a significant drop in performance between recommending friend links and recommending non-friend links.

algorithms. Learning information diffusion thus helps when it comes to building better SCF algorithms.

- Spectral Co-preference didn’t do well on friend link recommendations, however it did better on the non-friend link recommendations. When it comes to recommending friend links, friend interaction information coming through social regularizer seems more important than implicit co-likes information provided by the co-preference regularizer. When there is no social interaction information such as with non-friend links, co-preference methods with its implicit co-likes information appear much better than just vanilla collaborative filtering at projecting users into subspaces of common interest.

5.2.2 Summary

We summarize the observations made during the second trial:

- The social spectral regularization methods generally performed better in the live user trials, even when the results were split between friend link recommendations and non-friend link recommendations.
- Learning information diffusion models helps in SCF, as evidenced by the strong performance of Social Hybrid when recommending friend links.
- When there is no social interaction information, learning implicit co-likes information is better than using plain CF methods.
- The better performance of the social spectral regularization methods in the live trials were not reflected in the offline experiments. Perhaps there is a better metric than MAP that correlates with human preferences.

5.3 Algorithm Independent Observations

5.3.1 Click evidence

5.3.2 Impact of Popularity

5.3.3 Impact of Interactions

5.3.4 Impact of Steve Jobs

5.3.5 Impact of Genre

6. CONCLUSIONS

6.1 Summary

In this paper, we evaluated existing algorithms and proposed new algorithms for social collaborative filtering via the task of link recommendation on Facebook. Importantly, we outlined three main deficiencies in existing social collaborative filtering (SCF) matrix factorization (MF) techniques and proposed novel techniques in Section 3 to address them:

- Non-feature-based user similarity:** We extended existing social regularization and *social spectral regularization* methods to incorporate *user features* to learn user-user similarities in the latent space.
- Model direct user-user information diffusion:** We defined a new hybrid SCF method where we *combined* the *collaborative filtering (CF) matrix factorization (MF) objective* used by Matchbox [17] with a *linear content-based filtering (CBF) objective* used to model direct user-user information diffusion in the social network.
- Restricted common interests:** We defined a new social co-preference regularization method that *learns from pairs of user preferences* over the same item to learn *user similarities in specific areas* — a contrast to previous methods that typically enforce global user similarity when regularizing.

Having evaluated existing baselines (with minor extensions) and then evaluating these new algorithms in Section 5 in live online user trials with over 100 Facebook App users and data for over 30,000 unique Facebook users, we summarize the main results of the paper:

- In the first user trial, SCF MF beats all other baseline CF methods evaluated including MF.

- In the second user trial, each of social regularization, direct modeling of information features, and co-preference regularization improve performance over the best baseline from the first trial.
- Click feedback is not always as useful as explicit like/dislike ratings.
- Most popular links are not the most liked ones — they may be most liked by the most people, but they are not well-liked on average by everyone.

6.2 Future Work

This work just represents the tip of the iceberg in different improvements that SCF can make over more traditional non-social CF methods. Here we identify a number of additional future extensions that can potentially further improve the proposed algorithms in this paper:

- One critical feature that would have been useful is including a *genre* feature in the links (e.g., indicating whether the link represented a blog, news, video, etc.) to provide a fine-grained model of which types of links that users prefers to receive. This additional information would have likely prevented a number of observed dislikes from users regarding specific genres of links that they categorically disliked.
- Enforcing diversity in the recommended links would prevent redundant links about the same topic being recommended again and again. This is especially useful when an unusual event happens like the death of Steve Jobs and the ensuing massive amount of Steve Jobs related links that flooded Facebook. While users may like to see a few links on the topic, their interest in similar links decreases over time and diversity in recommendations could help address this saturation effect.
- Another future direction this work can go to is to incorporate active learning in the algorithms. This would ensure that the SCF algorithm did not exploit the learned preferences too much and made an active effort to discover better link preferences that are available.

While there are many exciting extensions of this work possible as outlined above, this paper represents a critical step forward in SCF algorithms based on top-performing MF methods and their ability to fully exploit the breadth of information available on social networks to achieve state-of-the-art link recommendation.

7. ADDITIONAL AUTHORS

Additional authors: Ehsan Abbasnejad (ANU & NICTA; Canberra, AU; email: first.last@nicta.com.au).

8. REFERENCES

- [1] M. Balabanović and Y. Shoham. Fab: content-based, collaborative recommendation. *Communications of the ACM*, 40:66–72, March 1997.
- [2] R. M. Bell and Y. Koren. Scalable collaborative filtering with jointly derived neighborhood interpolation weights. In *ICDM-07*, 2007.
- [3] J. Brown and P. Reinegen. Social ties and word-of-mouth referral behavior. *Journal of Consumer Research*, 1(3):350–362, 1987.
- [4] C.-C. Chang and C.-J. Lin. *LIBSVM: a Library for Support Vector Machines*, 2001.
- [5] C. Cortes and V. Vapnik. Support-vector networks. In *Machine Learning*, pages 273–297, 1995.
- [6] P. Cui, F. Wang, S. Liu, M. Ou, and S. Yang. Who should share what? item-level social influence prediction for users and posts ranking. In *International ACM SIGIR Conference (SIGIR)*, 2011.
- [7] K. Lang. NewsWeeder: Learning to filter netnews. In *12th International Conference on Machine Learning ICML-95*, pages 331–339, 1995.
- [8] W.-J. Li and D.-Y. Yeung. Relation regularized matrix factorization. In *IJCAI-09*, 2009.
- [9] D. C. Liu and J. Nocedal. On the limited memory BFGS method for large scale optimization. *Mathematical Programming*, 45(1):503–528, Aug 1989.
- [10] H. Ma, I. King, and M. R. Lyu. Learning to recommend with social trust ensemble. In *SIGIR-09*, 2009.
- [11] H. Ma, H. Yang, M. R. Lyu, and I. King. Sorec: Social recommendation using probabilistic matrix factorization. In *CIKM-08*, 2008.
- [12] H. Ma, D. Zhou, C. Liu, M. R. Lyu, and I. King. Recommender systems with social regularization. In *WSDM-11*, 2011.
- [13] A. Ng, M. Jordan, and Y. Weiss. On spectral clustering: Analysis and an algorithm. In *Advances in Neural Information Processing Systems NIPS 14*, 2001.
- [14] K. B. Petersen and M. S. Pedersen. *The matrix cookbook*, 2008.
- [15] P. Resnick and H. R. Varian. Recommender systems. *Communications of the ACM*, 40:56–58, March 1997.
- [16] R. Salakhutdinov and A. Mnih. Probabilistic matrix factorization. In *Advances in Neural Information Processing Systems*, volume 20, 2008.
- [17] D. H. Stern, R. Herbrich, and T. Graepel. Matchbox: large scale online bayesian recommendations. In *WWW-09*, pages 111–120, 2009.
- [18] S. H. Yang, B. Long, A. Smola, N. Sadagopan, Z. Zheng, and H. Zha. Like like alike: Joint friendship and interest propagation in social networks. In *WWW-11*, 2011.

APPENDIX

A. GRADIENT-BASED OPTIMIZATION

We seek to optimize sums of the objectives in Section 3 and will use gradient descent for this purpose.

For the overall objective, the partial derivative w.r.t. parameters \mathbf{a} are as follows:

$$\frac{\partial}{\partial \mathbf{a}} Obj = \frac{\partial}{\partial \mathbf{a}} \sum_i \lambda_i Obj_i = \sum_i \lambda_i \frac{\partial}{\partial \mathbf{a}} Obj_i \quad (16)$$

Anywhere a sigmoidal transform occurs $\sigma(o[\cdot])$, we can easily calculate the partial derivatives as follows

$$\frac{\partial}{\partial \mathbf{a}} \sigma(o[\cdot]) = \sigma(o[\cdot])(1 - \sigma(o[\cdot])) \frac{\partial}{\partial \mathbf{a}} o[\cdot]. \quad (17)$$

Hence anytime a $[\sigma(o[\cdot])]$ is optionally introduced in place of $o[\cdot]$, we simply insert $[\sigma(o[\cdot])(1-\sigma(o[\cdot]))]$ in the corresponding derivatives below.

Because most objectives below are not convex in U , V , or \mathbf{w} , we apply an *alternating gradient descent* approach [16]. In short, we take derivatives of U , V , and \mathbf{w} in turn while holding the others constant. Then we apply gradient descent in a round-robin fashion until we've reached local minima for all parameters; for gradient descent on one of U , V , or \mathbf{w} with the others held constant, we apply the L-BFGS optimizer [9] with derivatives defined below.

Before we proceed to our objective gradients, we define abbreviations for two useful vectors:

$$\begin{aligned}\mathbf{s} &= U\mathbf{x} & \mathbf{s}_k &= (U\mathbf{x})_k; \quad k = 1 \dots K \\ \mathbf{t} &= V\mathbf{y} & \mathbf{t}_k &= (V\mathbf{y})_k; \quad k = 1 \dots K \\ \mathbf{r} &= U\mathbf{z} & \mathbf{r}_k &= (U\mathbf{z})_k; \quad k = 1 \dots K\end{aligned}$$

All matrix derivatives used for the objectives below can be verified in [14].

$$\begin{aligned}\frac{\partial}{\partial U} Obj_{pmcf} &= \frac{\partial}{\partial U} \sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} \left(\underbrace{(R_{\mathbf{x}, \mathbf{y}} - [\sigma] \overbrace{x^T U^T V \mathbf{y}}^{o_{\mathbf{x}, \mathbf{y}}})}_{\delta_{\mathbf{x}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{y}) \in D} \delta_{\mathbf{x}, \mathbf{y}} [\sigma(o_{\mathbf{x}, \mathbf{y}})(1 - \sigma(o_{\mathbf{x}, \mathbf{y}}))] \mathbf{t} \mathbf{x}^T\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial V} Obj_{pmcf} &= \frac{\partial}{\partial V} \sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} \left(\underbrace{(R_{\mathbf{x}, \mathbf{y}} - [\sigma] \overbrace{x^T U^T V \mathbf{y}}^{o_{\mathbf{x}, \mathbf{y}}})}_{\delta_{\mathbf{x}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{y}) \in D} \delta_{\mathbf{x}, \mathbf{y}} [\sigma(o_{\mathbf{x}, \mathbf{y}})(1 - \sigma(o_{\mathbf{x}, \mathbf{y}}))] \mathbf{s} \mathbf{y}^T\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial U} Obj_{ru} &= \frac{\partial}{\partial U} \frac{1}{2} \text{tr}(U^T U) = U & \frac{\partial}{\partial V} Obj_{rv} &= V \\ \frac{\partial}{\partial \mathbf{w}} Obj_{rw} &= \frac{\partial}{\partial \mathbf{w}} \frac{1}{2} \mathbf{w}^T \mathbf{w} = \mathbf{w}\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial U} Obj_{rs} &= \frac{\partial}{\partial U} \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \frac{1}{2} \left(\underbrace{S_{\mathbf{x}, \mathbf{z}} - \mathbf{x}^T U^T U \mathbf{z}}_{\delta_{\mathbf{x}, \mathbf{y}}} \right)^2 \\ &= - \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \delta_{\mathbf{x}, \mathbf{y}} U(\mathbf{x} \mathbf{z}^T + \mathbf{z} \mathbf{x}^T)\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial U} Obj_{rss} &= \frac{\partial}{\partial U} \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} \frac{1}{2} S_{\mathbf{x}, \mathbf{z}}^+ (\mathbf{x} - \mathbf{z})^T U^T U (\mathbf{x} - \mathbf{z}) \\ &= \sum_{\mathbf{x}} \sum_{\mathbf{z} \in \text{friends}(\mathbf{x})} S_{\mathbf{x}, \mathbf{z}}^+ U (\mathbf{x} - \mathbf{z}) (\mathbf{x} - \mathbf{z})^T\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial \mathbf{w}} Obj_{phy} &= \frac{\partial}{\partial \mathbf{w}} \sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} \left(\underbrace{R_{\mathbf{x}, \mathbf{y}} - [\sigma] \overbrace{\mathbf{w}^T \mathbf{f}_{\mathbf{x}, \mathbf{y}} - [\sigma] \mathbf{x}^T U^T V \mathbf{y}}^{o_{\mathbf{x}, \mathbf{y}}^1}}_{\delta_{\mathbf{x}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{y}) \in D} \delta_{\mathbf{x}, \mathbf{y}} [\sigma(o_{\mathbf{x}, \mathbf{y}}^1)(1 - \sigma(o_{\mathbf{x}, \mathbf{y}}^1))] \mathbf{f}_{\mathbf{x}, \mathbf{y}}\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial U} Obj_{phy} &= \frac{\partial}{\partial U} \sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} \left(\underbrace{R_{\mathbf{x}, \mathbf{y}} - [\sigma] \mathbf{w}^T \mathbf{f}_{\mathbf{x}, \mathbf{y}} - [\sigma] \overbrace{\mathbf{x}^T U^T V \mathbf{y}}^{o_{\mathbf{x}, \mathbf{y}}^2}}_{\delta_{\mathbf{x}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{y}) \in D} \delta_{\mathbf{x}, \mathbf{y}} [\sigma(o_{\mathbf{x}, \mathbf{y}}^2)(1 - \sigma(o_{\mathbf{x}, \mathbf{y}}^2))] \mathbf{t} \mathbf{x}^T\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial V} Obj_{phy} &= \frac{\partial}{\partial V} \sum_{(\mathbf{x}, \mathbf{y}) \in D} \frac{1}{2} \left(\underbrace{R_{\mathbf{x}, \mathbf{y}} - [\sigma] \mathbf{w}^T \mathbf{f}_{\mathbf{x}, \mathbf{y}} - [\sigma] \overbrace{\mathbf{x}^T U^T V \mathbf{y}}^{o_{\mathbf{x}, \mathbf{y}}^2}}_{\delta_{\mathbf{x}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{y}) \in D} \delta_{\mathbf{x}, \mathbf{y}} [\sigma(o_{\mathbf{x}, \mathbf{y}}^2)(1 - \sigma(o_{\mathbf{x}, \mathbf{y}}^2))] \mathbf{s} \mathbf{y}^T\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial U} Obj_{rsc} &= \frac{\partial}{\partial U} \sum_{(\mathbf{x}, \mathbf{z}, \mathbf{y}) \in C} \frac{1}{2} \left(\underbrace{P_{\mathbf{x}, \mathbf{z}, \mathbf{y}} - \mathbf{x}^T U^T \text{diag}(V \mathbf{y}) U \mathbf{z}}_{\delta_{\mathbf{x}, \mathbf{z}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{z}, \mathbf{y}) \in C} \delta_{\mathbf{x}, \mathbf{z}, \mathbf{y}} \text{diag}(V \mathbf{y}) U (\mathbf{x} \mathbf{z}^T + \mathbf{z} \mathbf{x}^T)\end{aligned}$$

In the following, \circ is the Hadamard elementwise product:

$$\begin{aligned}\frac{\partial}{\partial V} Obj_{rsc} &= \frac{\partial}{\partial V} \sum_{(\mathbf{x}, \mathbf{z}, \mathbf{y}) \in C} \frac{1}{2} (P_{\mathbf{x}, \mathbf{z}, \mathbf{y}} - \mathbf{x}^T U^T \text{diag}(V \mathbf{y}) U \mathbf{z})^2 \\ &= \frac{\partial}{\partial V} \sum_{(\mathbf{x}, \mathbf{z}, \mathbf{y}) \in C} \frac{1}{2} \left(\underbrace{P_{\mathbf{x}, \mathbf{z}, \mathbf{y}} - (\overbrace{U \mathbf{x}}^{\mathbf{s}} \circ \overbrace{U \mathbf{z}}^{\mathbf{r}})^T V \mathbf{y}}_{\delta_{\mathbf{x}, \mathbf{z}, \mathbf{y}}} \right)^2 \\ &= - \sum_{(\mathbf{x}, \mathbf{z}, \mathbf{y}) \in C} \delta_{\mathbf{x}, \mathbf{z}, \mathbf{y}} (\mathbf{s} \circ \mathbf{r}) \mathbf{y}^T\end{aligned}$$